

# An Interpretable Retrieval-Augmented Generation (RAG) System with Local Knowledge Graph Explanations

*Technical Report*

## Team FPS

Pratyaksha Jha   Somita Agarwal   Freya Sheth  
B.Tech, Indian Institute of Technology Guwahati

January 12, 2026

## Abstract

Traditional Retrieval-Augmented Generation (RAG) systems focus on improving answer accuracy by grounding language models in retrieved documents, but they often fail to provide transparency regarding why a particular answer was generated. This limits trust, verifiability, and usability in high-stakes or academic settings. By utilizing a Local Contextual Entity Graph, the system provides a transparent pipeline that not only answers queries but also visualizes the metrics, entities, and evidence used to derive that answer. The system is designed to be client-side and privacy-preserving making sure that the data never leaves the user's machine. The system prioritizes clarity, traceability, and grounding over complex reasoning or large-scale optimization, aligning closely with the objectives of this challenge.

## Contents

---

<b>1</b>	<b>System Overview</b>	<b>3</b>
<b>2</b>	<b>Query Processing and Retrieval</b>	<b>3</b>
2.1	Query Processing and Retrieval . . . . .	3
<b>3</b>	<b>Entity Extraction and Graph Construction</b>	<b>3</b>
3.1	Entity Extraction . . . . .	3
3.2	Relationship Construction . . . . .	4
3.3	Graph Representation . . . . .	4
<b>4</b>	<b>Answer Generation</b>	<b>4</b>
4.1	Context Construction . . . . .	4
4.2	Generation Strategy . . . . .	4
4.3	The "I Don't Know" State . . . . .	4
<b>5</b>	<b>Preliminary Results &amp; Experiments</b>	<b>5</b>
5.1	Test Case: Query about jaundice from a book of Biochemistry . . . . .	5
<b>6</b>	<b>Interpretability and Transparency</b>	<b>5</b>
<b>7</b>	<b>Prototype</b>	<b>7</b>
7.1	Implementation Details . . . . .	7
7.2	Tech Stack . . . . .	7
7.3	Demo Example . . . . .	7
<b>8</b>	<b>Future Improvements</b>	<b>7</b>
<b>9</b>	<b>GitHub Repository:</b>	<b>7</b>
<b>10</b>	<b>Conclusion</b>	<b>8</b>

## System Overview

---

The pipeline follows a modular design centered on client-side processing to ensure data privacy and immediate interpretability.

1. **Document Ingestion:** Supports PDF, DOCX, and JSON using `pdf.js` and `mammoth.js`. Documents are fragmented into overlapping chunks to preserve local context.
2. **Retrieval Engine:** A term-frequency scoring algorithm with a specific **Factual Boost** for numerical data.
3. **Local Entity Extraction:** Pattern-based NLP to identify metrics and named entities within the retrieved subset.
4. **Extractive Synthesis:** A grounding-first approach that generates answers by selecting and ranking the most relevant source sentences.
5. **Explanation Layer:** A structured output consisting of color-coded entity tags and interactive evidence blocks.

Each component is intentionally designed to be modular, interpretable, and easy to reason about.

## Query Processing and Retrieval

---

### Query Processing and Retrieval

The system employs a Term-Frequency weighted retrieval method to identify the most relevant context for a given user query. The relevance score  $S$  for a document  $D$  given a query  $Q$  is calculated using the following objective function:

$$S(Q, D) = \sum_{t \in Q} (\text{count}(t, D) \times w_t) + 1_{\text{fact}}(D) \quad (1)$$

Where:

- $w_t$  is the term weight (configured to 2 for significant non-stopword terms).
- $\text{count}(t, D)$  represents the frequency of term  $t$  within document  $D$ .
- $1_{\text{fact}}(D)$  is a **Factual Boost** (+1 score) applied if the document contains numerical metrics or units (e.g., %, °C, currency), prioritizing data-rich chunks.

**Design Choice:** Keyword-based search was selected over dense vector embeddings for this prototype to ensure **exact-match transparency**. In explainable systems, it is critical that users can audit the retrieval process by identifying the specific lexemes that triggered a document's selection, thereby eliminating the "black box" nature of latent semantic spaces.

## Entity Extraction and Graph Construction

---

### Entity Extraction

Entities are extracted using high-performance Regular Expressions (Regex) to maintain system speed and transparency.

- **Technique:** Pattern-based Extraction.

- **Categories:**

- **METRIC:** Identified via patterns of numbers followed by units like %, °C, million, or billion.
- **ENTITY:** Identified by Title Case patterns (e.g., "Paris Agreement") for sequences of 1-3 capitalized words.

### Graph Logic:

- **Nodes:** The extracted Entities and Metrics.
- **Edges:** Implicitly defined by co-occurrence within the same retrieved text chunk. This creates a "Local Context Graph" that represents the specific world-view of the retrieved evidence, avoiding "hallucinated" connections from global pre-training.

## Relationship Construction

A local knowledge graph is constructed by mapping relationships between the query, extracted entities, and specific sentences:

- **Sentence-Level Relevance:** Sentences are scored based on the density of query terms and extracted entities.
- **Evidence Mapping:** The top 3 sentences per retrieved document are linked to the answer as "Eyes-on" evidence.

## Graph Representation

The graph is represented in the UI as a set of interactive badges and evidence blocks:

- **Nodes:** {text, type: METRIC | ENTITY}
- **Edges:** Implicit links between entities and their source sentences, presented as highlighted evidence blocks.

## Answer Generation

---

### Context Construction

The answer generation context consists of retrieved text chunks and highlighted sentences containing key entities.

### Generation Strategy

To meet the requirement of strict interpretability, the system utilizes **Extractive Synthesis** rather than abstractive generation. The system:

1. Scans retrieved chunks for sentences containing query keywords and extracted entities.
2. Ranks these sentences based on a local relevance score.
3. Concatenates the top 4 sentences to form a factual response.

### The "I Don't Know" State

A key feature of our design is the "No Answer Found" state. If no direct overlap exists between the query and the knowledge base, the system refuses to generate a response. This is a critical safety mechanism for transparent RAG, preventing the fabrications common in standard language models.

## Preliminary Results & Experiments

The system outputs a structured explanation alongside the final answer.

### Test Case: Query about jaundice from a book of Biochemistry

- **Query:** "jaundice"
- **Retrieval:** UCB levels in the blood become elevated (unconjugated hyperbilirubinemia), causing jaundice (Fig. Obstructive (posthepatic): In this instance, jaundice is not caused by overproduction of bilirubin or decreased conjugation but, instead, results from obstruction of the common bile duct (extrahepatic cholestasis). For example, the presence of a tumor or bile stones may block the duct, preventing passage. (Score: 19).
- **Evidence from Retrieved Documents:** Sentences and files from which entities and relations were derived.

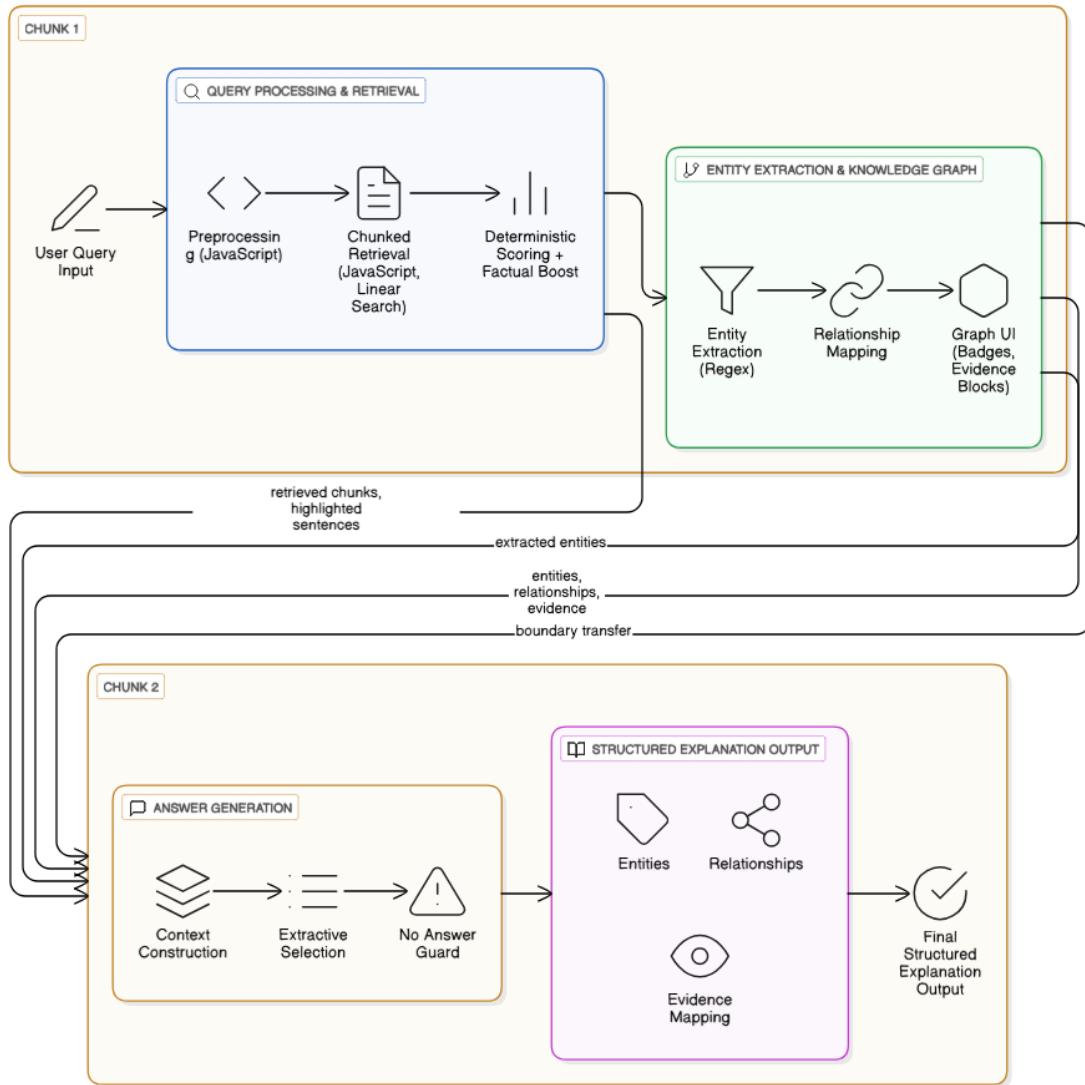
The screenshot shows the Explainable RAG System interface. At the top, there's a header with the title 'Explainable RAG System' and a sub-header 'Transparent retrieval-augmented generation with structured explanations'. Below this, it says '1296 documents in knowledge base'. There are buttons for '+ Add Documents' and a link icon. A 'Knowledge Base' section lists six documents from 'Lippincott's Illustrated Reviews Biochemistry, Second Edition' (chunks 1-6), each with its file name, chunk number, and character count. Below this is a search bar containing the query 'jaundice' and a 'Search' button. The main area is divided into two tabs: 'Answer & Evidence' (selected) and 'Source Documents'. Under 'Answer & Evidence', the generated answer is displayed: 'UCB levels in the blood become elevated (unconjugated hyperbilirubinemia), causing jaundice (Fig. Obstructive (posthepatic): In this instance, jaundice is not caused by overproduction of bilirubin or decreased conjugation but, instead, results from obstruction of the common bile duct (extrahepatic cholestasis). For example, the presence of a tumor or bile stones may block the duct, preventing passage of CB Figure 26.11 Jaundiced patient with the sclerae of his eyes)'.

Figure 1: Sample input and output.

## Interpretability and Transparency

Interpretability is achieved through:

- **Evidence Highlighting:** The system displays the exact sentences from the source documents that contributed to the score.
- **Structured Entity Tags:** Entities are color-coded (Blue for Metrics, Green for Entities) so users can see the "building blocks" of the information at a glance.
- **Relevance Scoring:** Each source document is displayed with its calculated score, allowing users to evaluate why one document was preferred over another.



**Figure 2: System Overview Architecture.** A comprehensive visualization of the four-stage Explainable RAG pipeline, mapping the flow from Query Processing to the final Structured Explanation Output.

## Prototype

---

### Implementation Details

The prototype is a browser-native application that requires no installation. It features a Knowledge Base management system, a real-time search interface, and a tabbed "Answer & Evidence" viewer.

### Tech Stack

- **Frontend UI:** HTML5, Tailwind CSS
- **Document Parsing:** `pdf.js` (PDF), `mammoth.browser.js` (DOCX)
- **Processing Logic:** Vanilla JavaScript (ES6+)
- **Retrieval/Graph:** Regex-based extraction and term-frequency scoring
- **Deployment:** Static Web App (Client-side execution)

### Demo Example

**Query:** "What is the economic impact of climate change?"

**Output:** "The economic cost of climate inaction is estimated at \$23 trillion by 2050. Green bonds have reached \$500 billion in issuance. Some economists argue the transition costs are underestimated.. Global temperatures have risen by 1.2°C since pre-industrial times."

### Future Improvements

---

While the current prototype establishes a baseline for transparency, the following enhancements are proposed to increase the depth of reasoning and scalability:

- **Multimodal Data Ingestion (OCR):** Integrating Optical Character Recognition (OCR) engines like *Tesseract* or *PaddleOCR* to process scanned documents and images. This would allow the system to extract text from diagrams and flowcharts, which often contain critical process-related information.
- **Structured Table Parsing:** Implementing layout-aware extraction (using tools like `Unstructured.io` or `Camelot`) to convert tables into structured formats like Markdown or JSON. This ensures that the relational data within cells is preserved, allowing the RAG system to answer quantitative queries that require row-column lookups.
- **Hybrid Retrieval:** Integration of semantic vector embeddings (e.g., *SBERT*) alongside the current keyword-based search to capture latent context while maintaining exact-match transparency.
- **SVO Triplet Extraction:** Moving beyond simple co-occurrence to Subject-Verb-Object (SVO) extraction to transform the local graph into a more actionable relational Knowledge Graph.

### GitHub Repository:

---

The source code for this prototype is hosted at the following link:

**Repo link:** [Retriever](#)

## Conclusion

---

This work presents an end-to-end Explainable Retrieval-Augmented Generation (RAG) system that integrates lightweight knowledge graph construction to improve transparency and trust. The X-RAG prototype demonstrates that explainability does not require complex graph databases. By focusing on Local Contextual Extraction and Transparent Scoring, we provide a system where the user can verify every claim. This approach reduces the "hallucination" risk inherent in standard RAG systems and provides a clear audit trail for sensitive data analysis.